Oviva: a safe first-line assistant that gives coaches time back

Prepared by Vinod Kumar G for Oviva, October 2026

What I understood

Oviva raised a USD 220M Series D to scale personalised care, on an AI-enabled platform that lets clinicians focus on patients [1]. Your AI stack is Python, Vertex AI, Gemini and the Agent Development Kit [2]. Your GP page lists 8,000 NHS patients treated monthly [3]. The app manual asks patients not to send urgent questions through chat, because their health care professional may not check it every day [4]. In Germany, from 2026, 20% of a DiGA's payment depends on measured successful use [5].

The opportunity

Many patient messages are routine: how the programme works, eligibility, app how-tos. An assistant can answer those at once from approved content, with the page shown. That only helps if it is provably safe: anything clinical, urgent or account-specific must reach a person, with context, every time. And that safety needs a number before each release, not a demo.

What it could be worth

Roughly £36,000 to £256,000 a year in UK coach time alone. That excludes any effect on engagement and renewals.

- Formula: patients a month × routine messages per patient × share answered automatically × minutes saved × 12 × cost per hour.
- Low: 8,000 × 1 × 25% × 3 min = 1,200 hours a year at £30. High: 8,000 × 2 × 40% × 5 min = 6,400 hours a year at £40.
- Every input except the 8,000 patients is my assumption, to be replaced by your numbers.

What I built to test the idea

safe-support (https://github.com/VinodKumarGgithub/safe-support) is an assistant over Oviva's public UK pages. It applies safety rules first, then searches the approved content, and a release check scores every change. On 64 labelled questions:

- It recognised 10 of 10 urgent messages and answered 0 clinical questions automatically.
- It answered all 30 routine questions, 28 of them from the right page.
- Its own check blocked the first version at 79% right-page answers. A ranking fix raised that to 93% without loosening a safety rule. Gemini can sit behind the same rules and check.

A pilot

Four weeks in shadow mode, on one market and one question type (programme and app questions). You share anonymised questions and approved answers. The assistant drafts replies for staff, and no patient sees them. We would measure the share of drafts accepted without edits, the hand-overs that reached the right team, the time to first reply, and zero unsafe answers in review.

Why me

- I lead a 12-member team on a healthcare (Pharmacy Benefit Management) platform handling 1M+ transactions a day for 11 enterprise clients, at 99.98% uptime.
- I own core business logic that turns complex healthcare rules into maintainable Node.js/TypeScript services.
- I built Lara, a personal multi-agent RAG assistant with LangGraph agents, tool calling and vector search, traced and evaluated in LangSmith and Langfuse.

Sources

1. https://www.startupticker.ch/en/news/digitalhealth-company-oviva-banks-usd-200m-series-d-funding
2. https://oviva.com/global/en/engineering/
3. https://oviva.com/uk/en/refer/
4. https://oviva.com/uk/en/user-manual/
5. https://www.pharmazeutische-zeitung.de/hersteller-muessen-erfolg-ihrer-diga-kuenftig-belegen-152588/

Vinod Kumar G · imvinodkumarg@gmail.com · imvinod.com · github.com/VinodKumarGgithub